

Homework 2

Hanlin He (hxx160630)

October 8, 2018

Chapter 7

Exercise 85

Consider a thread acquired and release the lock for the first time. Let *qnodeX* be the node created for the thread. After the thread released the lock, we will have:

- *tail = qnodeX*
- *qnodeX.locked = false*

If the same thread tries to acquire the lock again, the execution of `lock` method would be:

```
1 public void lock() {  
2     Qnode qnode = myNode.get();           // qnode ← qnodeX  
3     qnode.locked = true;                   // qnodeX.locked ← true  
4     Qnode pred = tail.getAndSet(qnode);   // pred ← qnodeX  
5     while (pred.locked) {}                // get stuck in pred.locked forever  
6 }
```

At line 4, return tail from `tail.getAndSet(qnode)` will return *qnodeX* itself. And the *locked* value was already set to *true*, so the **while** loop will never break.

Exercise 91

- `testAndSet()` spin lock:

```
1 public boolean isLocked() {  
2     return state.get();  
3 }
```

- CLH queue lock:

```
1 public boolean isLocked() {  
2     return tail.get().locked;  
3 }
```

- MCS queue lock:

```

1 public boolean isLocked() {
2     return tail.get() != null;
3 }

```

Chapter 8

Exercise 95

```

1 /**
2  * 95-0. Interface definition.
3  */
4 public interface Account {
5     void deposit(int k);
6     void withdraw(int k);
7     int getBalance();
8 }

```

```

1 import java.util.concurrent.locks.Condition;
2 import java.util.concurrent.locks.Lock;
3 import java.util.concurrent.locks.ReentrantLock;
4
5 /**
6  * 95-1. Implement this savings account using locks and conditions.
7  */
8 public class SavingsAccount implements Account {
9
10     private int balance;
11
12     final Lock mutex = new ReentrantLock();
13
14     private final Condition canWithdraw = this.mutex.newCondition();
15
16     public SavingsAccount() {
17         this(0);
18     }
19
20     public SavingsAccount(int balance) {
21         this.balance = balance;
22     }
23
24     /**
25      * Return current balance.
26      */
27     @Override
28     public int getBalance() {
29         this.mutex.lock();
30         try {
31             return this.balance;
32         } finally {
33             this.mutex.unlock();
34         }
35     }
36 }

```

```

34     }
35 }
36
37 /**
38  * deposit(k) adds k to the balance
39  */
40 @Override
41 public void deposit(int k) {
42     this.mutex.lock();
43     try {
44         this.balance += k;
45         this.canWithdraw.signalAll();
46     } finally {
47         this.mutex.unlock();
48     }
49 }
50
51 /**
52  * withdraw(k) subtracts k from balance, if the balance is at least
53  * k,
54  * and otherwise blocks until the balance becomes k or greater
55  */
56 @Override
57 public void withdraw(int k) {
58     this.mutex.lock();
59     try {
60         while (k > this.balance) {
61             this.canWithdraw.awaitUninterruptibly();
62         }
63         this.balance -= k;
64     } finally {
65         this.mutex.unlock();
66     }
67 }

```

```

1  import java.util.concurrent.locks.Condition;
2
3  /**
4   * 95-2. Implement two kinds of withdrawals for SavingsAccount:
5   * ordinary and preferred.
6   */
7  public class SavingsAccountWithPreferredWithdraw extends
8      SavingsAccount {
9
10     private final Condition canOrdinaryWithdraw =
11         this.mutex.newCondition();
12     private int preferredWithdrawCount = 0;
13
14     /**
15      * ordinaryWithdraw perform withdrawal that yield to preferred
16      * withdrawal.
17      */
18     public void ordinaryWithdraw(int k) {
19         // if preferred withdrawals waiting, wait for them to finish
20         this.mutex.lock();

```

```

17     try {
18         while (this.preferredWithdrawCount > 0) {
19             this.canOrdinaryWithdraw.awaitUninterruptibly();
20         }
21         // if no preferred withdrawals waiting, perform withdraw
22         this.withdraw(k);
23     } finally {
24         this.mutex.unlock();
25     }
26 }
27
28 /**
29  * preferredWithdraw perform withdrawal with higher priority.
30  */
31 public void preferredWithdraw(int k) {
32     // increment preferred withdraw count using mutex
33     this.mutex.lock();
34     try {
35         this.preferredWithdrawCount++;
36     } finally {
37         this.mutex.unlock();
38     }
39
40     // perform withdraw
41     this.withdraw(k);
42
43     // decrement preferred withdraw count using mutex
44     this.mutex.lock();
45     try {
46         this.preferredWithdrawCount--;
47         this.canOrdinaryWithdraw.signalAll();
48     } finally {
49         this.mutex.unlock();
50     }
51 }
52 }

```

```

1 import java.util.concurrent.locks.Lock;
2 import java.util.concurrent.locks.ReentrantLock;
3
4 /**
5  * 95-3. Add a transfer() method that transfers a sum from one account
6  * to another.
7  */
8 public class SavingsAccountWithTransfer extends SavingsAccount {
9
10     private final Lock transferMutex = new ReentrantLock();
11
12     /**
13      * transfer() method transfers a sum from one account to another
14      */
15     public void transfer(int k, Account reserve) {
16         this.transferMutex.lock();
17         try {
18             reserve.withdraw(k);
19             this.deposit(k);

```

```

19     } finally {
20         this.transferMutex.unlock();
21     }
22 }
23
24 /*
25  If there are less than 10 threads for each account, all transfer
26  methods called at 1:00 are guaranteed to return at 2:00.
27  */
28 }

```

Chapter 8

Exercise 24

For the first history

Let H_1 denotes the history, α denotes the operation $A : r.read(1)$, β and γ denote the operations $B : r.write(1)$ and $B : r.read(2)$, and δ denotes the operation $C : r.write(2)$:

- Quiescently Consistent: Yes.

After execution of α , β and γ , the program became quiescent. Thus, the history would be quiescently consistent if the result of δ appear after α , β and γ . Suppose γ was the last operation to take effect among α , β and γ , the history would be legal and quiescently consistent.

- Sequentially Consistent: Yes.

Consider the sequential history shown in eq. (1).

$$S_1 \equiv \beta \rightarrow \alpha \rightarrow \delta \rightarrow \gamma \quad (1)$$

S_1 is apparently equivalent with H_1 . And the object subhistory of r is legal. Thus, the history H_1 is sequentially consistent.

- Linearizable: Yes.

Consider the sequential history in eq. (1) again. The preceding relations of H_1 is

$$\rightarrow_{H_1} \equiv \{\alpha \rightarrow \gamma, \beta \rightarrow \gamma, \delta \rightarrow \gamma, \} \quad (2)$$

Apparently, $\rightarrow_{H_1} \subseteq \rightarrow_{S_1}$. Thus, the history H is linearizable.

For the second history

Let H_2 denotes the history, α denotes the operation $A : r.read(1)$, β and γ denote the operations $B : r.write(1)$ and $B : r.read(1)$, and δ denotes the operation $C : r.write(2)$:

- Quiescently Consistent: Yes.

After execution of α , β and γ , the program became quiescent. Thus, the history would be quiescently consistent if the result of δ appear after α , β and γ . Suppose β was the last operation to take effect among α , β and γ , the history would be legal and quiescently consistent.

- Sequentially Consistent: Yes.

Consider the sequential history S_2 shown in eq. (3).

$$S_2 \equiv \delta \rightarrow \beta \rightarrow \alpha \rightarrow \gamma \quad (3)$$

S_2 is apparently equivalent with H_2 . And the object subhistory of r is legal. Thus, the history H_2 is sequentially consistent.

- Linearizable: Yes.

Consider the sequential history in eq. (3) again. The preceding relations of H_2 is

$$\rightarrow_{H_2} \equiv \{\alpha \rightarrow \gamma, \beta \rightarrow \gamma, \delta \rightarrow \gamma, \} \quad (4)$$

Apparently, $\rightarrow_{H_2} \subseteq \rightarrow_{S_2}$. Thus, the history H_2 is linearizable.

Exercise 27

Since Java does not guarantee linearizability, the definition of buffer would lead to linearizability problems.

```
1 T[] items = (T[]) new Object[Integer.MAX_VALUE];
```

The `items` was defined as an array of generic type `T`, instead of an array of `AtomicReference<T>` or simply `AtomicReferenceArray<T>`. The changes of one element in `items` within one thread `items[slot] = x;` may not be noticeable to other threads immediately. So it is possible to have two threads T_1 and T_2 with a history H that:

$$H : T_1.enq(1) \rightarrow T_2.deq() \text{ throw empty exception} \quad (5)$$

because `value = items[slot]` in T_2 was still `null` after $T_1.enq(1)$. However, the legal sequence history for eq. (5) is:

$$S : T_2.deq() \text{ throw empty exception} \rightarrow T_1.enq(1) \quad (6)$$

Apparently, $\rightarrow_H \not\subseteq \rightarrow_S$, the implementation is not linearizable.